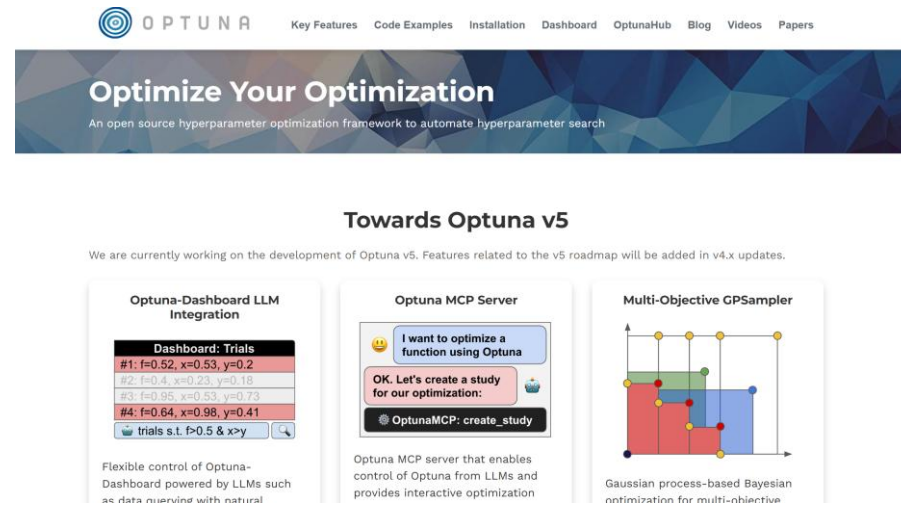


# Hyperparameter optimization using Optuna



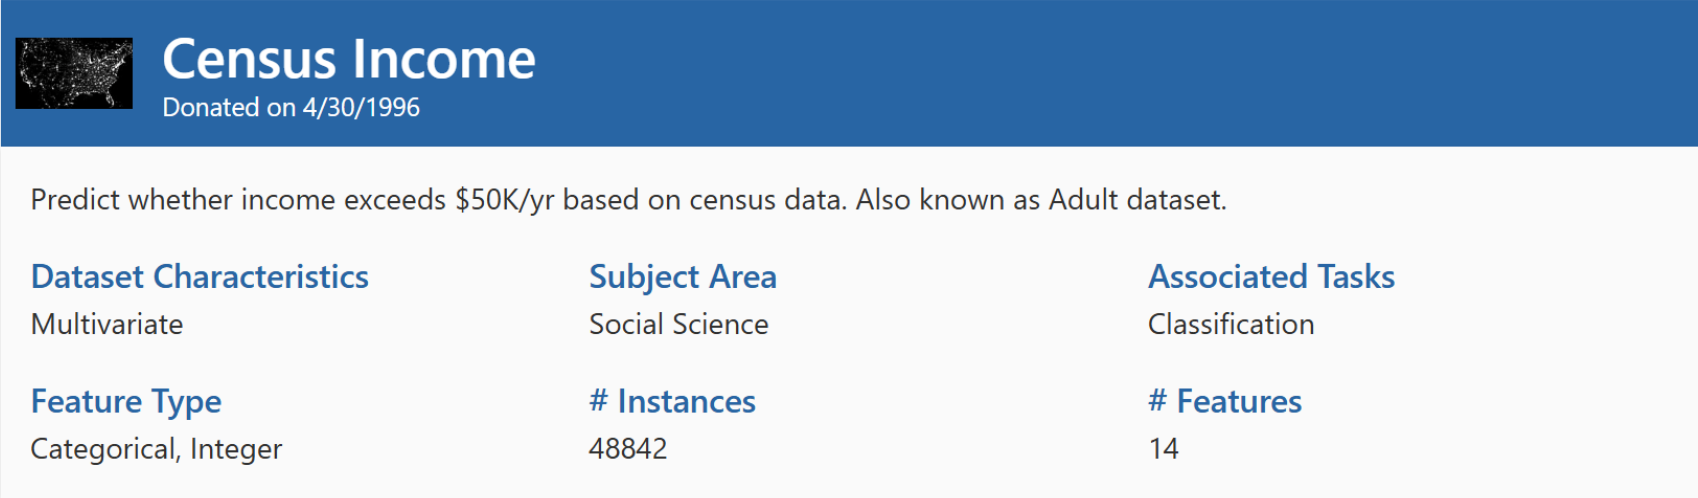
Clinical AI Study Group  
Amit Saha


# Hyperparameter tuning

- ML algorithms have model parameters and objective functions – which are optimized during the learning process
- Hyperparameter tuning is the process of intelligently searching the search space of the parameters that are specified for the learning – i.e., user input/configuration

# Today's demo

- XGBoost (eXtreme Gradient Boosting) as the ML algorithm for classification
  - Ensemble method, multiple trees

- e income



## Census Income

Donated on 4/30/1996

Predict whether income exceeds \$50K/yr based on census data. Also known as Adult dataset.

Dataset Characteristics	Subject Area	Associated Tasks
Multivariate	Social Science	Classification
Feature Type	# Instances	# Features
Categorical, Integer	48842	14

# XGBoost

$$\text{Final Prediction} = \sum_{m=1}^M \eta \cdot f_m(x)$$

- $M$ : number of estimators (trees)
- $f_m(x)$ : the  $m$ -th tree's prediction
- $\eta$ : learning rate
- each  $f_m$  is a small regression tree (a "weak learner")

# XGBoost – Hyperparameters of interest

- Learning rate: step size of model weight update
- Number of estimators: number of trees
- Max depth: depth of each tree

## Guidelines

1. Small learning rate with larger number of estimators
  1. Balance between stability and training speed (lower learning rate)
2. More depth results in potential overfitting the training data
3. Higher number of estimators (boosting rounds), better accuracy

# Hyperparameter tuning using Optuna

- ML algorithm agnostic
- Study – set up range of the parameters, which sampling (optimization algorithm) algorithm to use, and whether to maximize or minimize objective(s)
- Trial – a single run of the ML algorithm + the parameters chosen from the study configuration

# Objective function – single objective

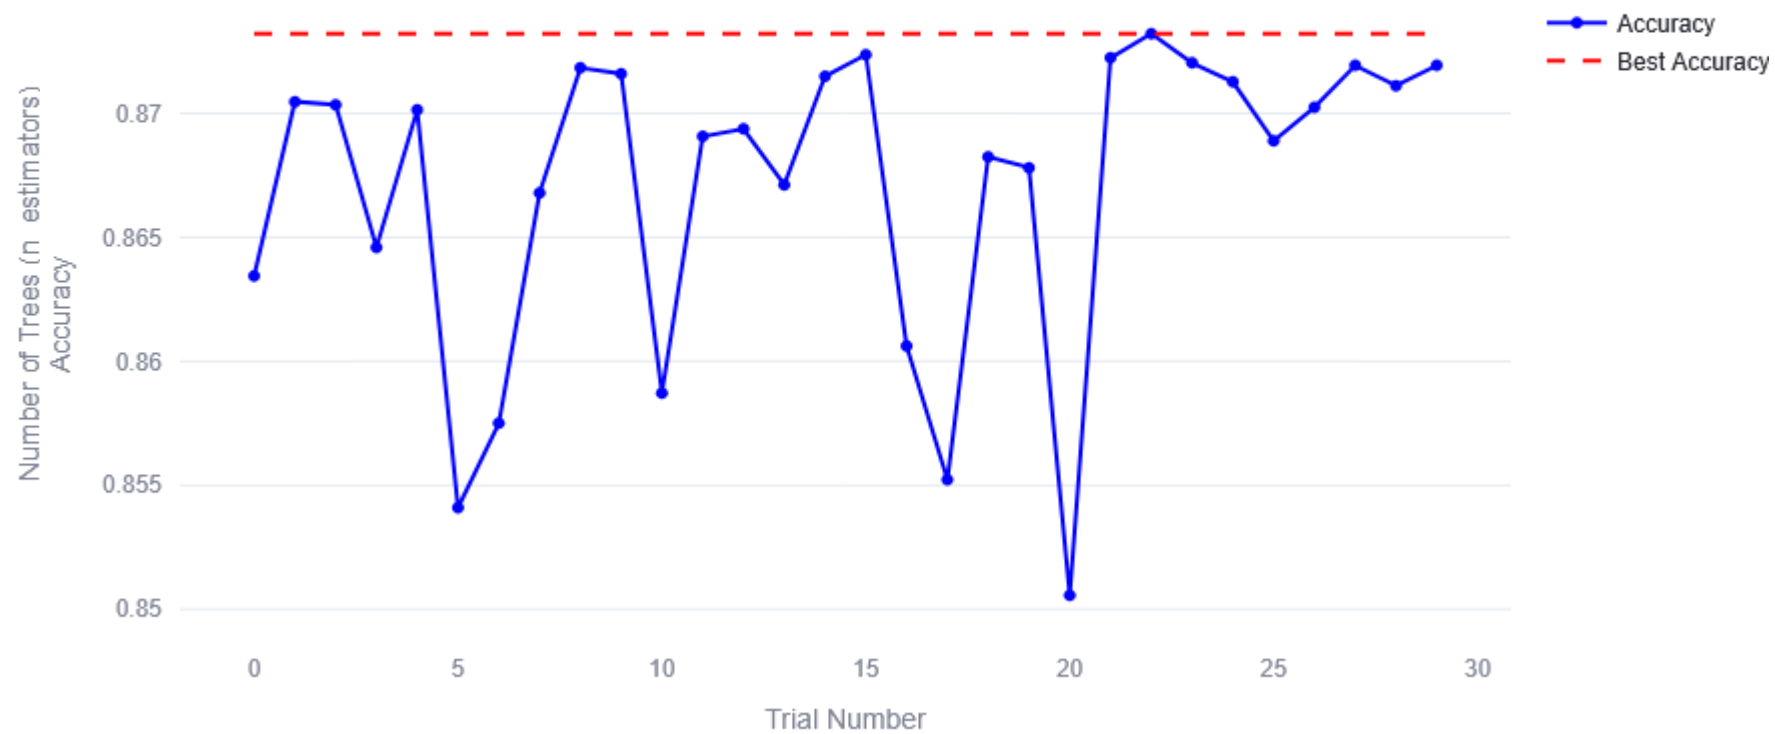
```
def objective(trial, X_train, y_train):  
    param = {  
        'max_depth': trial.suggest_int('max_depth', 3, 10),  
        'learning_rate': trial.suggest_float('learning_rate', 0.01, 0.3, log=True),  
        'n_estimators': trial.suggest_int('n_estimators', 50, 300)  
    }  
    model = XGBClassifier(**param)  
  
    scores = cross_val_score(model, X_train, y_train, cv=3, scoring='accuracy', n_jobs=-1, error_score='raise')  
  
    return float(scores.mean())
```

# Study definition – single objective

```
def run_single_objective_optimization(n_trials=50):  
  
    study = optuna.create_study(direction='maximize', study_name='xgboost_single_objective')  
  
    study.optimize(lambda trial: objective(trial, X_train, y_train), n_trials=n_trials)  
  
    best_params = study.best_params.copy()  
    best_params.update({  
        'random_state': 42,  
        'eval_metric': 'logloss',  
    })  
  
    final_model = XGBClassifier(**best_params)  
    final_model.fit(X_train, y_train)  
  
    # Evaluate on test set  
    test_accuracy = final_model.score(X_test, y_test)  
  
    return study
```



## L<sub>1</sub> Optimization History



Learning Rate vs Number of Trees



DEMO

# Objective function – multi objective

```
def objective(trial, X_train, y_train):
    # Define hyperparameter search space
    param = {
        'max_depth': trial.suggest_int('max_depth', 3, 10),
        'learning_rate': trial.suggest_float('learning_rate', 0.01, 0.3, log=True),
        'n_estimators': trial.suggest_int('n_estimators', 50, 300),
    }

    # Create XGBoost classifier
    model = XGBClassifier(**param)

    # Objective 1:
    auc_scores = cross_val_score(
        model, X_train, y_train, cv=3, scoring='roc_auc', n_jobs=-1, error_score='raise'
    )
    val_auc = auc_scores.mean()
    performance = val_auc

    complexity = param['n_estimators'] * param['max_depth']


    return performance, complexity
```

# Study definition – multi objective

```
def run_multi_objective_optimization(n_trials=50):  
    study = optuna.create_study(  
        directions=['maximize', 'minimize'], # maximize roc auc, minimize complexity  
        study_name='xgboost_multi_objective'  
    )  
  
    study.optimize(lambda trial: objective(trial, X_train, y_train), n_trials=n_trials)  
  
    # Pareto optimal solutions  
    study.best_trials
```

DEMO

## Pareto Frontier: ROC-AUC vs Model Complexity



	RandomSampler	GridSampler	TPESampler	CmaEsSampler	NSGAIISampler	QMCSampler	GPSampler	BoTorchSampler
Float parameters	✓	✓	✓	✓	▲	✓	✓	✓
Integer parameters	✓	✓	✓	✓	▲	✓	✓	▲
Categorical parameters	✓	✓	✓	▲	✓	▲	✓	▲
Pruning	✓	✓	✓	▲	✗ (▲ for single-objective)	✓	▲	▲
Multivariate optimization	▲	▲	✓	✓	▲	▲	✓	✓
Conditional search space	✓	▲	✓	▲	▲	▲	▲	▲
Multi-objective optimization	✓	✓	✓	✗	✓ (▲ for single-objective)	✓	✗	✓
Batch optimization	✓	✓	✓	✓	✓	✓	▲	✓
Distributed optimization	✓	✓	✓	✓	✓	✓	▲	✓
Constrained optimization	✗	✗	✓	✗	✓	✗	✗	✓
Time complexity (per trial) (*)	$O(d)$	$O(dn)$	$O(dn \log n)$	$O(d^3)$	$O(mp^2)$ (***)	$O(dn)$	$O(n^3)$	$O(n^3)$
Recommended budget (# trials) (****)				1000 - 10000	100 - 10000	as many as one likes	- 500	10 - 100

- Optuna calls them [samplers](#)

## Optimization algorithms used

- Single objective – [TPESampler](#) (Tree-structured Parzen Estimator)
- Multi-objective – [NSGA-IISampler](#) (Evolutionary algorithm based optimizer)
- There is also an option to choose the sampling algorithm “automatically” based on your data - [AutoSampler](#)



# Conclusion

- Optuna makes hyperparameter tuning more manageable rather than having to manually try out various ranges
- Slides, and demos: <https://github.com/amitsaha/optuna-demo>